Emerald Feedback Wall

Product Requirements Document

Version: 0.1 — Draft
Status: For Review
Owner: TBD
Last updated: March 2026

---

1. Overview

The Emerald Feedback Wall is a conference-deployed, large-format interactive kiosk that collects structured B2B client feedback across existing product screens and forward-looking AI/agentic concepts. All feedback is synthesised overnight via AI into a prioritised roadmap and overnight development brief, which is then used to rapidly prototype changes before Day 2 of the event.

The product serves two purposes simultaneously:

- **Research instrument** — capturing high-quality, structured product feedback from real clients in a low-friction, contextually-rich environment
- **Marketing moment** — demonstrating to clients that their input directly and visibly shapes the product, with a Day 2 reveal proving speed-to-build

---

2. Problem Statement

Traditional conference feedback mechanisms (surveys, workshops, booth conversations) produce data that feeds into a quarterly roadmap cycle at best. Clients rarely see evidence that their feedback changed anything. Meanwhile, product teams attending conferences have rare, concentrated access to a cross-section of real users but no structured way to capture and act on what they hear.

The Emerald Feedback Wall compresses the feedback-to-prototype cycle from months to hours, making client-centricity visible and tangible in real time.

---

3. Goals & Success Metrics

Goal	Metric	Target
Capture high-quality feedback	Responses per event	≥ 50 responses / day
Drive Day 2 engagement	Attendees who recall feedback being actioned	≥ 70% (post-event survey)
Generate marketing assets	Consent-approved quotes captured	≥ 10 per event
Enable overnight sprint	AI synthesis → prototype brief in < 15 mins	✓
Prove product velocity	Day 2 prototype demo	✓ shipped

---

4. Users & Roles

4.1 Attendees (Primary — Day 1)

B2B clients and prospects attending the conference. They interact with the feedback wall to give input on existing screens and react to AI/agentic product concepts. Assumed to be non-technical business users from operations, product, finance, and engineering roles.

4.2 Emerald Product Team (Internal — Day 1 Evening)

Uses the Synthesis tab to generate the prioritised roadmap and PRD. Drives overnight prototyping sprint based on the output. Needs clean, actionable output with minimal post-processing.

4.3 Emerald Marketing Team (Internal — Day 2)

Uses consent-approved quotes and the Day 2 reveal moment as a live proof point of product velocity. Needs attribution-safe, publish-ready content.

4.4 Conference Facilitator (Internal — Both Days)

Manages the kiosk physically, monitors response volume, and triggers synthesis at the designated time. Needs a simple admin view and the ability to lock the synthesis panel from attendee access.

---

5. Scope

5.1 In Scope — v1.0 (Conference MVP)

Left Drawer — Three-tab panel

- **Features tab** — Feature request cards with upvoting. Add new requests inline. Sorted by vote count. Workflow context field per request.
- **Kudos tab** — Freetext quote submission with role attribution and explicit public consent checkbox. Renders submitted quotes live.
- **Synthesis tab** — Admin-gated. Displays live signal summary (vote totals, screen feedback count, kudos count). Toggle between Roadmap and PRD output modes. Streams AI-generated output. Copy-to-clipboard export.

Hero — Screen Feedback (always visible)

- Level 1: Application selector (Core Platform, Agentic Workflows, Analytics, AI Concepts)
- Level 2: Screen grid with representative wireframe thumbnails per app
- Level 3: Enlarged screen mock + feedback panel (type tag + freetext). Submission saves to synthesis data pool.
- AI Concepts section visually differentiated (dark treatment) to signal forward-looking content vs. current product

Global

- Live response counter in top bar
- Synthesis countdown timer
- Collapsible drawer for max hero space
- Responsive to large-format display (min 1440px wide recommended)

5.2 Out of Scope — v1.0

- Card sort module (deprioritised; can be added as a fifth left-drawer tab in v1.1)
- Real-time projector/secondary display output
- Attendee login or identity capture
- Integration with existing roadmap tools (Jira, Linear, Productboard)
- Post-event data export pipeline
- Multi-event data persistence

6. Functional Requirements

6.1 Feature Requests

- FR-01 Users can submit a feature request via an inline text input (no page navigation)
- FR-02 Users can upvote existing feature requests (single tap; no account required)
- FR-03 Feature cards display vote count and are sorted descending by votes
- FR-04 Each feature request optionally captures a "workflow context" field explaining the use case
- FR-05 New submissions appear at the top of the list before vote sorting

6.2 Kudos & Quotes

- KD-01 Users can submit a freetext quote of any length
- KD-02 Users can optionally select a role (Ops, Eng, Product, Finance, Exec)
- KD-03 An explicit consent checkbox must be checked before the quote can be marked as public-safe
- KD-04 Public-safe quotes are visually tagged ("✓ Public OK") in the drawer
- KD-05 Submitted kudos render live in the drawer for all users to see

6.3 Screen Feedback

- SF-01 Navigating to a screen does not require selecting a feedback type first
- SF-02 Screen thumbnails are representative wireframe mocks (not live product screenshots in v1)
- SF-03 Feedback type selection (Pain point / Confusing / Missing element / Works well / Suggestion) is required before submission
- SF-04 Freetext field is optional but encouraged via contextual placeholder copy
- SF-05 Submitted feedback is tagged with app, screen name, and type for synthesis grouping
- SF-06 Post-submission, the user is offered a prompt to feedback another screen

6.4 AI Synthesis

- SY-01 Synthesis tab is access-controlled via a PIN (4–6 digit, set by conference facilitator)
- SY-02 Signal summary shows live counts: total feature votes, screen feedback items, kudos submitted
- SY-03 Two output modes: **Roadmap** (P0/P1/P2 + patterns + marketing quotes) and **PRD** (problem statement, user stories, acceptance criteria, design guidance)
- SY-04 Output streams token-by-token for perceived liveness

- SY-05 Output renders structured markdown (headers, bullets, numbered lists)
- SY-06 Copy-to-clipboard button exports full output as plain text
- SY-07 Synthesis prompt includes: all feature requests with vote counts and workflow context, all screen feedback with type tags and app/screen attribution, all kudos with role labels, and counts of total signals
- SY-08 Roadmap output must identify P0 items as prototypable within 8 hours
- SY-09 PRD output must include at minimum: overview, problem statement, in-scope / out-of-scope, user stories, acceptance criteria, design guidance
- SY-10 Synthesis runs against Anthropic Claude via API (commercial terms — no data training by default)

6.5 Privacy & Data Handling

- PV-01 All AI synthesis calls use the Anthropic API under Commercial Terms of Service (not consumer Claude Pro/Team)
- PV-02 No personally identifiable information is collected at any point
- PV-03 Role attribution on kudos is the maximum granularity of identity captured
- PV-04 Consent for public use of quotes must be explicitly captured at point of submission; default is private
- PV-05 A privacy notice ("Your response is processed securely. No names are stored.") is displayed on the kudos submission form
- PV-06 Data is not persisted beyond the event without a separate data retention decision

---

7. Non-Functional Requirements

Area	Requirement
Display	Optimised for large-format touchscreen or monitor (≥ 1440px). Mouse/touch compatible.
Performance	Screen navigation < 200ms. AI synthesis first token < 3s.
Availability	Offline-capable for core capture functionality (synthesis requires connectivity)
Security	Synthesis tab PIN-gated. No external auth dependencies for attendee flows.
Accessibility	Touch targets ≥ 44px. Sufficient colour contrast. No time-limited interactions.

---

8. Information Architecture

```
Feedback Wall
├── Left Drawer (collapsible)
│   ├── Features Tab
│   │   ├── Add new (inline)
│   │   └── Feature cards (upvotable, sorted by votes)
│   ├── Kudos Tab
│   │   ├── Add new (inline, with role + consent)
│   │   └── Submitted quotes
│   └── Synthesis Tab [PIN-gated]
│       ├── Signal summary
│       ├── Mode toggle (Roadmap / PRD)
│       ├── Generate button
│       └── Streaming output + copy
└── Hero — Screen Feedback
    ├── Level 1: App selector
    │   ├── Core Platform
    │   ├── Agentic Workflows
    │   ├── Analytics
    │   └── AI Concepts ✦ (dark treatment)
    ├── Level 2: Screen grid (thumbnails)
    └── Level 3: Screen detail + feedback panel
        ├── Screen mock (enlarged)
        ├── Feedback type selector
        ├── Freetext input
        └── Submit → confirmation → next screen prompt
```

---

9. Design Principles

Always-on hero. Screen feedback is the primary research instrument and should never be hidden behind navigation. The drawer is supplementary.

Visible progress. The response counter and synthesis countdown make the overnight sprint feel like a live event, not a passive survey. This is part of the marketing value.

Friction-proportionate depth. Upvoting an existing feature request takes one tap. Adding a new one requires two steps. Full screen feedback requires three. Depth scales with commitment.

Consent by design. Public use of quotes requires an explicit, visible opt-in at the point of submission. The default is always private.

Prototype framing. Screen mocks should always be labelled "Representative wireframe · not final UI" to set expectations and avoid premature debates about visual polish.

Synthesis as a handoff, not a decision. The AI output is a starting point for the overnight team, not a commitment. The PRD and Roadmap are drafts to react to, not ship verbatim.

---

10. AI Synthesis Prompt Design

Roadmap Mode Output Structure

1. **P0 — Build Tonight** (1–2 items, 8-hour constraint explicit)
2. **P1 — Next Sprint** (3–4 items with rationale)
3. **P2 — Backlog** (remaining themes)
4. **Patterns & Insights** (2–3 cross-cutting observations)
5. **Marketing Moments** (2–3 verbatim consent-approved quotes with role)

PRD Mode Output Structure

1. **Overview** (one paragraph)
2. **Problem Statement** (evidence-based from feedback)
3. **Scope — Tonight's Build** (ruthlessly scoped)
4. **Out of Scope**
5. **User Stories** (5–7, role / action / outcome format)
6. **Acceptance Criteria** (per P0 feature)
7. **Design Guidance** (UX principles from feedback themes)
8. **Success Metrics** (Day 2 demo success definition)

Prompt Inputs (assembled dynamically)

- Feature requests: text + vote count + workflow context
- Screen feedback: app + screen name + type tag + freetext
- Kudos: text + role (consent-approved only surfaced as quotes)
- Metadata: total signal counts per category

---

11. Open Questions

#	Question	Owner	Priority
1	What are the real application names and screens to populate for the first event?	Product	P0
2	Should AI Concepts screens show actual concept mocks or remain as wireframe placeholders?	Design	P0
3	What PIN management process works for conference facilitators?	Engineering	P1
4	Should a secondary "admin view" display aggregated feedback in real time (e.g. on a backstage monitor)?	Product	P1
5	What is the data handling policy after the event — delete, archive, or import to roadmap tool?	Legal / Product	P1
6	Should the Day 2 prototype be revealed on the same kiosk, or a separate screen?	Marketing	P1
7	Is card sort needed for v1, or moved to v1.1?	Product	P2
8	Do we want response counts broken down by category (e.g. 12 on Agentic Workflows, 8 on Analytics) to guide synthesis weighting?	Engineering	P2

12. Risks & Mitigations

Risk	Likelihood	Impact	Mitigation
Feedback volume too low for meaningful synthesis	Medium	High	Staff the kiosk; brief facilitators to actively direct attendees; set synthesis threshold at 30 responses minimum
AI synthesis hallucinates patterns not in data	Medium	Medium	Human review step before overnight team acts on output; synthesis clearly framed as a draft
Prototype looks finished; clients expect shipping	High	Medium	Badge all screens "Prototype · Non-prod · Synthetic data"; verbal framing in Day 2 reveal
Sensitive client info captured in freetext	Low	High	No PII fields; privacy notice on submission; synthesis prompt instructs model to generalise, not quote individuals

Risk	Likelihood	Impact	Mitigation
Synthesis API unavailable overnight	Low	High	Export raw feedback data to CSV as fallback; manual synthesis protocol as backup
Public kudos quotes used without clear consent	Low	High	Consent checkbox is required field; only consent-flagged quotes surfaced in synthesis marketing section

13. Milestones

Milestone	Description	Target
M1	Wireframe sign-off	Week 1
M2	Real screen assets and app names confirmed	Week 2
M3	Functional prototype with real API connection	Week 3
M4	PIN gating, copy export, privacy notice	Week 3
M5	Facilitator run-through and kiosk display testing	Week 4
M6	Conference Day 1 deployment	Event
M7	Overnight synthesis + prototype sprint	Event night
M8	Day 2 reveal	Event
M9	Post-event retrospective + v1.1 scope	Week after